

└ Introduction

NOTE: To whom ever is reading this document. Pandoc processes divs that inherit from the 'notes' class as a note. These will not show up in your slides. To compile slides with notes run `make notes`. You will need to be running pandoc `-version >= 1.12.A`

- All of the examples in these lectures will make use of C++11 features. C++11 provides thread abstractions that should in theory work on any OS.

Concurrent Programming

└ Problem

Problem

- Why would we want to use threads?
- Let's say we have a service that performs a simple task. Take in sets of number sequences and sorts them using quicksort.
 $f(S) = S'$ such that
 $S = \{s_i | s_i = \{a_1, \dots, a_n\}, 1 \leq i \leq m, m, n \in \mathbb{N}\}$ and
 $S' = \{S' | s_i \text{ is sorted}\}$. $f(S)$ will then call $Q(s_i)$ on each sequence m times.

- Quicksorting sets of subsequences is a good example since the subsequences can be of any size and the quicksort's performance may vary.

Concurrent Programming

Serial Example

Serial Example

- A simple program may contain a single loop in f that calls the quicksort function at each iteration.

```
int main() {
    // A vector with 10,000 vectors of size 1,000 with unsorted integers.
    std::vector<std::vector<int>> v(1000);
    // ...
    std::vector<int> i(1000);

    for(std::vector<std::vector<int>> &v : v) {
        // ...
        sort(v.begin(), v.end());
    }
}

void sort(std::vector<int> &v) { // ... }
```

- Code is in Concurrency/parallel_soring_example.

Concurrent Programming

└ Parallel Example

Parallel Example

- A program that utilizes threads to speed up the process would spawn multiple threads.

```
int main() {
    // A vector with 10,000 vectors of size 1,000 with unsorted integers.
    std::vector<std::vector<int>>> v = { std::vector<int>(1000),
                                        // ...
                                        std::vector<int>(1000) };

    std::vector<std::thread> threads;
    // Open thread to sort each subvector.
    for(std::vector<std::vector<int>>>::iterator a_i = v.begin();
        a_i != v.end(); ++a_i)
    {
        threads.push_back(std::thread(sort, std::ref(*a_i)));
    }

    // Wait for each thread to complete its work.
    for(std::vector<std::thread>::iterator t = threads.begin();
        t != threads.end(); ++t)
    {
        (*t).join(); // or t->join();
    }

    void sort(std::vector<int> & to_sort) { /* ... */ }
```

- Spawning 10,000 threads probably isn't a smart idea.

Concurrent Programming

└ Use Cases

Use Cases

- Operating Systems: Linux, Windows, and Unix.
- Graphical interfaces use event driven multithreading to preserve responsiveness.
- Games, separation of input, physics, and rendering.
- Web server technologies such as databases, search engines, and web servers.
- HMMER
- Bioinformatics.

- http://www.valvesoftware.com/publications/2007/GDC2007_SourceMulticore.pdf
- Web server technologies such as Apache, Nginx, and Microsoft ISS employ multithreading to separate requests.
- http://en.wikipedia.org/wiki/Event-driven_programming

2015-01-08

Concurrent Programming

└ History

└ Early Multithreading & Multitasking Systems

- Early Machines
 - Single process model.
 - Batch processing.
- Berkeley Timesharing System
 - Give processes **time-slots** of execution.
 - Memory is shared.
 - Computer remains usable for other operators.
- Unix
 - Processes now have dedicated memory.
 - Later, threading support added. Subprocesses that share memory with the process.

- `http://www.faqs.org/faqs/os-research/part1/section-10.html` provides a nice and simple intro to the history of threading.
- Chapter 4 in *Operating System Concepts* contains additional information about threads.
- `http://en.wikipedia.org/wiki/Unix`

Concurrent Programming

└ Threading Architecture

└ Threading Models

Threading Models

- kernel threads (most kernel threads on Linux are processes).
- user threads (threads that processes spawn).




Figure 1: Threading Models. Retrieved from Operating System Concepts

- Refer to Chapter 4: **Operating System Concepts 8e.** for more information on threading models.
 - Book states that Linux uses the one-to-one model.

Concurrent Programming

└ Threading Architecture

└ Hardware

Hardware

- Duplication of registers to store multiple states (Intel Hyper-threading).
 - Threads can still lock while waiting for CPU resources.
- Intel Hyper-Threading.
 - Multiple cores.
 - Multiple sockets.



Classic Design

Intel®



- Intel Hyper-Threading Image adopted from <https://software.intel.com/en-us/articles/performance-insights-to-intel-hyper-threading-technology?language=es>.
- Multi-core processor diagram was retrieved from https://en.wikipedia.org/wiki/Multi-core_processor

Concurrent Programming

└ Multithreaded Programming

└ Minimal Working Example

Minimal Working Example

```
void sort(std::vector<int> & ta_sort) {
    std::sort(ta_sort.begin(), ta_sort.end());
}

void parallel_sorting(std::vector<std::vector<int>> & V) {
    std::vector<std::thread> threads;
    for(std::vector<std::vector<int>> & v : itasortar v, i = V.begin();
        v, i != V.end(); ++i) {
        threads.push_back(std::thread(sort, std::ref(*v, i)));
    }

    for(std::vector<std::thread> & threads : threads) {
        t = threads.end(); ++t;
    }
    (*t).join();
}

int main(int argc, char * argv[]) {
    std::vector<std::vector<int>> V;
    fill_with_vectors(V, 10000, 1000);
    parallel_sorting(V);
    return 0;
}
```

- Includes at the top of the example.
 - `stdlib.h`, `iostream`, `algorithm`, `thread`, and `utils.h`.
- `utils.h` contains `fill_with_vectors(std::vector< std::vector<int> >, int, int)`.

Concurrent Programming

└ Multithreaded Programming

└ Threads in C++11

- ▶ After we have created T with vectors of random integers. We pass T to the `parallel_sorting` function.
- ▶ The first line of code `std::vector< std::thread > threads` is a vector that contains thread objects. If thread needs to be terminated with later then it is necessary to keep track of it.
- ▶ Second we iterate through the vectors in T and spawn a thread to sort each vector with:
 - `threads.push_back(std::thread(sort, std::ref(*v_i)));`
- ▶ `sort` is the function that defines the instructions for the thread instance.
- ▶ `std::ref(*v_i)` tells the thread constructor that `sort` requires `std::vector<int> &` as a parameter.

- `#include<thread>` a C++11 wrapper for various system threads. In the case of Linux and Unix it is a PThread (POSIX) thread wrapper.

Concurrent Programming

└ Multithreaded Programming

└ Defining a Thread Function

- Threads require a function to execute since they are subprocesses.
- The thread function `sort` is a wrapper for `std::sort` from the C++ Standard Library (`stdlib`)

```
void sort(std::vector< int > & to_sort) {  
    std::sort(to_sort.begin(), to_sort.end());  
}
```

- The STL is not the C++ Standard Library. It just stands for Standard Library which was created by Alexander Stepanov.
 - <http://stackoverflow.com/questions/5205491/whats-this-stl-vs-c-standard-library-fight-all-about>

Concurrent Programming

└ Thread Communication

└ Race Condition

Race Condition

- When `requests_served++` is executed a race condition may occur.
- `requests_served++` is not an atomic operation. Expanding it to machine code would result in:

```

register1 = requests_served    (1)
register1 = register1 + 1      (2)
requests_served = register1    (3)

```

- If a context change were to arise between lines 1 and 2 or 1 and 3. Then there is the possibility that `requests_served` will have changed due to another thread. Which would make `register1` inconsistent with `requests_served`.

- What problems may arise from the code?
 - Race conditions. *Critical Section*
 - Operations that appear atomic may not be once compiled.
 - Provide Assembly example of race condition.

Concurrent Programming

- └ Thread Communication
 - └ Critical Section

- Any section of code that reads or writes to data that is shared amongst threads.
- Must satisfy three requirements ensure consistency.
 1. **Mutual Exclusion:** If a thread is in a critical section then other threads must wait for it to exit the section.
 2. **Progress:** A thread cannot wait inside of a critical section. Waiting can cause a deadlock.
 3. **Bounded Waiting:** Threads shall not board the critical section.

- Refer to page 228 chapter 6 in *Operating System Concepts* 8th Edition.

Concurrent Programming

└ Thread Communication

└ Semaphores

- A mutex is a semaphore that only allows a single thread to access a critical section.
- More advanced data structure that provides mutual exclusion access to a critical section.
- Unlike a classic mutex a semaphore keeps count of the threads that want to access a resource.
- Designed to allow multiple threads access a critical section.
- Two operations used.
 - `wait(semaphore)`: Thread is blocked until another thread calls `signal`.
 - `signal(semaphore)`: Thread calls `signal` to indicate exit of critical section and allow another thread to enter.

- Explanation and definition of semaphores came from lecture 6 in the lecture materials from *CSE 120: Principles of Operating by Systems Alex C. Snoeren* (also included in the resources directory).
- Page 189 from the *Art of Multiprocessor Programming* by *Maurice Herlihy & Nir Shavit*.

Concurrent Programming

└ Thread Communication

└ Building a Semaphore in C++11

Building a Semaphore in C++11

```

Semaphore C++11

class semaphore {
private:
    std::mutex _mtx;
    std::condition_variable _cv;
    int _count;

public:
    semaphore(int count = 1): _count(count) {}
    void signal() {
        std::lock_guard<std::mutex> lck(_mtx);
        _count++;
        _cv.notify_one();
    }
    void wait() {
        std::unique_lock<std::mutex> lck(_mtx);
        // Over Dispatch (local) function w/
        _cv.wait(lck, [&lck]{ return _count > 0; });
        _count--;
    }
};
  
```

- It most cases it is recommended to use an existing library for semaphores.
- Such as boost or the POSIX defined `semaphore.h`
- `lock_guard` will lock within the scope using a mutex.
- `unique_lock` allows the condition variable to associate a set of threads to a common lock and defer their execution.
- `condition_variable`
 - `notify_one()` will wake up a sleeping thread.
 - `wait(unique_lock &, predicate)` will put a thread to sleep. predicate determines if a thread should go back to sleep after a spurious wakeup.
 - predicate is a C++11 anonymous function.

Concurrent Programming

└ Thread Communication

└ Readers-Writers Problem

- The semaphore can then be used to synchronize communication between a writer thread and set of reader threads.

Reader

```
void writer(shared_data & shared_data, semaphore & wrt) {
    while(1) {
        wrt.wait();
        //Access to shared data.
        wrt_signal();
        total_writes++;
    }
}
```

- Refer to page 241 from *Operating System Concepts 8 e.* or slide 10 in the lecture slides provided by Alex C. Snoeren CSE 120.
- Refer to `main.cpp` in the `semaphore_example` for a working demonstration of the readers/writers problem.

Concurrent Programming

└ Thread Communication

└ Readers-Writers Problem

Readers-Writers Problem

```

Writer
void reader(readdata & shared_data, semaphore & wrt,
           int & read_count) {
    while(1) {
        wrt.wait();
        read_count++;
        if(read_count == 1) {
            wrt.wait();
        }
        // Read the shared data.
        rd_signal();
        wrt.wait();
        read_count--;
        if(read_count == 0) {
            wrt.signal();
        }
    }
}

```

- Refer to page 241 from *Operating System Concepts 8 e.* or slide 10 in the lecture slides provided by Alex C. Snoeren CSE 120.
- Refer to `main.cpp` in the `semaphore_example` for a working demonstration of the readers/writers problem.


Concurrent Programming

Thread Communication

Deadlocks

Deadlocks

- What can cause deadlocks? There exists four conditions.
 - Mutual exclusion:** There exists at least a single resource that can be held in a non-sharable mode.
 - Hold and wait:** Thread holds onto a resource and waits for another resource to be freed.
 - No preemption:** Threads cannot be stripped of their resources by other threads.
 - Circular wait:** When two or more threads are holding and waiting on shared resources.



- Chapter 7 from *Operating System Concepts* covers deadlocks and deadlock avoidance.
 - Section 7.2 covers the four conditions necessary for a deadlock.
 - <http://nob.cs.ucdavis.edu/classes/ecs150-1999-02/dl-cond.html>
- Figure (d) Thread A goes into the critical zone and holds a resource and waits for another resource to be freed. Thread B starves while waiting.

Concurrent Programming

Thread Communication

Resource-Allocation Graph

Resource-Allocation Graph




Figure 2: Adopted from Operating System Concepts

- Chapter 7.4 Deadlock Prevention from *Operating System Concepts* and the Chapter 7 deadlock slides has more information.
- Represent resource usage as a directed acyclic graph.
 - If any cycles exist then there may exist a deadlock.

Concurrent Programming

└ Thread Communication

└ Lock-Free Programming

Lock-Free Programming

- ▶ Another method to prevent deadlocks is to use lock-free constructs.
- ▶ A lock-free structure guarantees throughput, but doesn't prevent starvation.
- ▶ Need to make use of atomic operations to construct the lock free data structure.
 - ▶ Atomic types, for example `std::atomic<T>`.
 - ▶ Atomic compare and swap (CAS). Such as `std::atomic_compare_exchange_*`

Advantages

- ▶ Guarantees no deadlocks.
- ▶ Scalable.

Disadvantages

- ▶ May be slower than lock-based structures.
- ▶ More difficult to implement.

- Page 60 from *The Art of Multiprocessor Programming* provides a nice definition of lock-free programming.
- Dr. Dobbs has a nice article about lock-free programming.
 - <http://www.drdobbs.com/lock-free-data-structures/184401865>.
- Lock-Free programming is a difficult subject.
- libcds provides a set of concurrent data structures
<http://libcds.sourceforge.net/>
- Nice set of Youtube videos about lock-free programming.
 - CppCon by Herb Sutter Part 1:
<https://www.youtube.com/watch?v=c1g09aB9nbs>
 - CppCon by Herb Sutter Part 2:
 -

Concurrent Programming

└ Inter-Process Communication

└ Forking

- How about process level parallelism? Use the `fork` command.
- `forking` a process will create a *child* (new) process that is an exact copy of the *parent* (calling) process except:
 - Locks are not preserved (including file locks).
 - Other threads from the parent process are not copied. Only the thread that forked the process is copied.
- Process IDs are not preserved. The child will be assigned new IDs.
- For more exceptions refer to the POSIX.1-2008 specifications for `fork()`.

- Provide simple code example and image.
- Section 4.4.1 in *Operating System Concepts* mentions that there exists fork implementations that will duplicate all threads when `fork` is called. Most versions of `fork` will only duplicate the thread that called the function. The book doesn't provide any references so it may be safe to assume that `fork` only duplicates the calling thread.
- `man 2 fork` on OS X and Linux will provide usage details.
- The Open Group provides the specifications of `fork` on a POSIX.1-2008 compatible system.
<http://pubs.opengroup.org/onlinepubs/9699919799/>

Concurrent Programming

└ Inter-Process Communication

└ Forking

Forking

- How many times will printf be called? How will we know what printf was called by the root process?

Forking Example

```
int main(int argc, char * argv[])
{
    pid_t pid0 = fork(); // fork returns 0 to the child process.
    pid_t pid1 = fork();
    printf("(%d, %d)\n", pid0, pid1); // print two unsigned numbers.
}
```


- fork will return the child pid to the parent and 0 to the child.
- printf requires #include <stdio.h>
- fork requires #include <unistd.h>
- Program Output:
(36961, 36962)
(36961, 0)
(0, 36963)
(0, 0)

Concurrent Programming

└ Inter-Process Communication

└ Forking

Forking
Forking Diagram



- A Visual example of the process behind `fork()`.
- The graph below shows the inheritance order.

Concurrent Programming

└ Inter-Process Communication

└ Pipes

Pipes

- Shared between the parent and child process (or multiple children). Different processes cannot share pipes.
- Enables communication between the parent and child process.
- Ordinary pipes provide unidirectional communication so that one end can be written to (write-end) and the other read from (read-end).

Pipe



Figure 3: Visual Representation of a Pipe: Adopted from Operating System Concepts.

- Provide simple code example and image.
- Refer to *Operating System Concepts* section 3.6.3: Pipes.
- Refer to the man pages
 - man 2 pipe
 - man 2 read
 - man 2 write

Concurrent Programming

└ Inter-Process Communication

└ Named Pipes: FIFO

- Referred to as First-In First-Out (FIFO) on POSIX/Unix systems.
- Enables communication between separate processes.
- Represented as a special file handle that points to a location in memory.
- Functionality similar to pipes, except bidirectional communication is possible. Unlike a pipe, reading and writing from the same file descriptor is possible.
- More overhead.
 1. Create the FIFO file.
 2. Open the FIFO.
 3. Read/Write to the FIFO.

- Provide simple code example and image.
- Refer to the operating systems book for more information.
- Refer to the man pages.

```
– man 2 mkfifo  
– man 2 open  
– man 2 read  
– man 2 write  
– man 2 close
```

Concurrent Programming

└ Inter-Process Communication

└ Sockets

Sockets

- Sockets provide full-duplex communication streams.
- Remote connections across the network.
- Primary tool to setup client-server communication model.

Advantages

- Dynamic: Allows the distribution of processes across multiple machines.
- More abstracted since the machines could be running their own operating system.

Disadvantages

- More overhead to set up.
 - Create a socket.
 - Bind the socket to an address.
 - Connect to the socket.
- Slower than FIFOs and Pipes since the data passes through the network stack.

- Refer to the man pages:
 - `man 2 socket`

Concurrent Programming

└ Threading Revisited

└ Schedulers

- Pools of threads do not guarantee optimal execution.
- Different threads will have varying execution times.
- Use a scheduler to ensure the desired optimal performance is achieved.
- Using a defined set of heuristics the scheduler will dequeue and run the desired threads from the pool.

- There isn't any defined thread pool class in C++11.

Concurrent Programming

└ Limitations

└ Amdahl's Law

Amdahl's Law

- Determine the potential code speedup with Amdahl's Law.
 - This version assumes that a case of parallelization.

$$T(N) = \frac{1}{(1-P) + \frac{P}{N}}$$

- P is a value between 0 and 1. P is the fraction of the program that is executed in parallel.
- As P approaches 1 then the program becomes more parallelized.
- If $P \rightarrow 1$ then the program is solving an embarrassingly parallel problem. The speedup is linear to the number of cores.
- N is the count of processors.
- Amdahl's law assumes a fixed problem size. Which causes a diminishing returns effect as the number of cores increase.

- http://en.wikipedia.org/wiki/Amdahl%27s_law
- <http://www.drdobbs.com/parallel/amdahls-law-vs-gustafson-barsis-law/240162980>
- Amdahl's Law assumes the dataset that the program is working on is static in size.
- For example a program that only parallelizes a fixed sized problem would follow this rule. Let's say that a program will always multiply two 1000x1000 matrices. Then the speedup of that program will follow Amdahl's law as the number of processors increase.

Concurrent Programming

└ Limitations

└ Gustafson's Law

- Unlike Amdahl's Law, Gustafson's Law assumes that the program will work on larger problems to utilize all of the processors.

$$S(N) = N - P(N - 1)$$

- P is the fraction of the program that is parallel.
- N is the number of processors.

- http://en.wikipedia.org/wiki/Gustafson%27s_law
- <http://www.drdobbs.com/parallel/amdahls-law-vs-gustafson-barsis-law/240162980>